

strucchange: An R Package for Testing for Structural Change in Linear Regression Models

Achim Zeileis¹ Friedrich Leisch¹ Kurt Hornik¹ Christian Kleiber²

¹ Institut für Statistik & Wahrscheinlichkeitstheorie, Technische Universität Wien,

² Institut für Wirtschafts- und Sozialstatistik, Universität Dortmund,

Abstract

This introduction to the R package **strucchange** is a (slightly) modified version of the paper of [Zeileis, Leisch, Hornik, and Kleiber \(2002\)](#), which reviews tests for structural change in linear regression models from the generalized fluctuation test framework as well as from the F test (Chow test) framework. It introduces a unified approach for implementing these tests and presents how these ideas have been realized in **strucchange**. Enhancing the standard significance test approach the package contains methods to fit, plot and test empirical fluctuation processes (like CUSUM, MOSUM and estimates-based processes) and to compute, plot and test sequences of F statistics with the $\sup F$, $\text{ave} F$ and $\text{exp} F$ test. Thus, it makes powerful tools available to display information about structural changes in regression relationships and to assess their significance. Furthermore, it is described how incoming data can be monitored.

Keywords: structural change, CUSUM, MOSUM, recursive estimates, moving estimates, monitoring, R, S.

1 Introduction

The problem of detecting structural changes in linear regression relationships has been an important topic in statistical and econometric research. The most important classes of tests on structural change are the tests from the generalized fluctuation test framework ([Kuan and Hornik 1995](#)) on the one hand and tests based on F statistics ([Hansen 1992a](#); [Andrews 1993](#); [Andrews and Ploberger 1994](#)) on the other. The first class includes in particular the CUSUM and MOSUM tests and the fluctuation test, while the Chow and the $\sup F$ test belong to the latter. A topic that gained more interest rather recently is to monitor structural change, i.e., to start after a history phase (without structural changes) to analyze new observations and to be able to detect a structural change as soon after its occurrence as possible.

This paper concerns ideas and methods for implementing generalized fluctuation tests as well as F tests in a comprehensive and flexible way, that reflects the common features of the testing procedures. It also offers facilities to display the results in various ways.

This paper is organized as follows: In Section 2 the standard linear regression model upon which all tests are based will be described and the testing problem will be specified. Section 3 introduces a data set which is also available in the package and which is used for the examples in this paper. The following sections 4, 5 and 6 will then explain the tests, how they are implemented in **strucchange** and give examples for each. Section 4 is concerned with computing empirical fluctuation processes, with plotting them and the corresponding boundaries and finally with testing for structural change based on these processes. Analogously, Section 5 introduces the F statistics and their plotting and testing methods before Section 6 extends the tools from Section 4 for the monitoring case.

2 The model

Consider the standard linear regression model

$$y_i = x_i^\top \beta_i + u_i \quad (i = 1, \dots, n), \quad (1)$$

where at time i , y_i is the observation of the dependent variable, $x_i = (1, x_{i2}, \dots, x_{ik})^\top$ is a $k \times 1$ vector of observations of the independent variables, with the first component equal to unity, u_i are iid($0, \sigma^2$), and β_i is the $k \times 1$ vector of regression coefficients. Tests on structural change are concerned with testing the null hypothesis of “no structural change”

$$H_0 : \quad \beta_i = \beta_0 \quad (i = 1, \dots, n) \quad (2)$$

against the alternative that the coefficient vector varies over time, with certain tests being more or less suitable (i.e., having good or poor power) for certain patterns of deviation from the null hypothesis.

It is assumed that the regressors are nonstochastic with $\|x_i\| = O(1)$ and that

$$\frac{1}{n} \sum_{i=1}^n x_i x_i^\top \longrightarrow Q \quad (3)$$

for some finite regular matrix Q . These are strict regularity conditions excluding trends in the data which are assumed for simplicity. For some tests these assumptions can be extended to dynamic models without changing the main properties of the tests; but as these details are not part of the focus of this work they are omitted here.

In what follows $\hat{\beta}^{(i,j)}$ is the ordinary least squares (OLS) estimate of the regression coefficients based on the observations $i + 1, \dots, i + j$, and $\hat{\beta}^{(i)} = \hat{\beta}^{(0,i)}$ is the OLS estimate based on all observations up to i . Hence $\hat{\beta}^{(n)}$ is the common OLS estimate in the linear regression model. Similarly $X^{(i)}$ is the regressor matrix based on all observations up to i . The OLS residuals are denoted as $\hat{u}_i = y_i - x_i^\top \hat{\beta}^{(n)}$ with the variance estimate $\hat{\sigma}^2 = \frac{1}{n-k} \sum_{i=1}^n \hat{u}_i^2$. Another type of residuals that are often used in tests on structural change are the recursive residuals

$$\tilde{u}_i = \frac{y_i - x_i^\top \hat{\beta}^{(i-1)}}{\sqrt{1 + x_i^\top (X^{(i-1)\top} X^{(i-1)})^{-1} x_i}} \quad (i = k + 1, \dots, n), \quad (4)$$

which have zero mean and variance σ^2 under the null hypothesis. The corresponding variance estimate is $\tilde{\sigma}^2 = \frac{1}{n-k} \sum_{i=k+1}^n (\tilde{u}_i - \bar{\tilde{u}})^2$.

3 The data

The data used for examples throughout this paper are macroeconomic time series from the USA. The data set contains the aggregated monthly personal income and personal consumption expenditures (in billion US dollars) between January 1959 and February 2001, which are seasonally adjusted at annual rates. It was originally taken from <http://www.econmagic.com/>, a web site for economic times series. Both time series are depicted in Figure 1.

The data is available in the **strucchange** package: it can be loaded and a suitable subset chosen by

```
> library("strucchange")
> data("USIncExp")
> USIncExp2 <- window(USIncExp, start = c(1985, 12))
```




Figure 1: Personal income and personal consumption expenditures in the US

We use a simple error correction model (ECM) for the consumption function similar to [Hansen \(1992b\)](#):

$$\Delta c_t = \beta_1 + \beta_2 e_{t-1} + \beta_3 \Delta i_t + u_t, \quad (5)$$

$$e_t = c_t - \alpha_1 - \alpha_2 i_t, \quad (6)$$

where c_t is the consumption expenditure and i_t the income. We estimate the cointegration equation (6) by OLS and use the residuals \hat{e}_t as regressors in equation (5), in which we will test for structural change. Thus, the dependent variable is the increase in expenditure and the regressors are the cointegration residuals and the increments of income (and a constant). To compute the cointegration residuals and set up the model equation we need the following steps in R:

```
> coint.res <- residuals(lm(expenditure ~ income, data = USIncExp2))
> coint.res <- lag(ts(coint.res, start = c(1985, 12), freq = 12),
+   k = -1)
> USIncExp2 <- cbind(USIncExp2, diff(USIncExp2), coint.res)
> USIncExp2 <- window(USIncExp2, start = c(1986, 1), end = c(2001,
+   2))
> colnames(USIncExp2) <- c("income", "expenditure", "diff.income",
+   "diff.expenditure", "coint.res")
> ecm.model <- diff.expenditure ~ coint.res + diff.income
```

Figure 2 shows the transformed time series necessary for estimation of equation (5).

In the following sections we will apply the methods introduced to test for structural change in this model.

4 Generalized fluctuation tests

The generalized fluctuation tests fit a model to the given data and derive an empirical process, that captures the fluctuation either in residuals or in estimates. For these empirical processes the limiting processes are known, so that boundaries can be computed, whose crossing probability under the null hypothesis is α . If the empirical process path crosses these boundaries, the fluctuation is improbably large and hence the null hypothesis should be rejected (at significance level α).




Figure 2: Time series used – first differences and cointegration residuals

4.1 Empirical fluctuation processes: function `efp`

Given a formula that describes a linear regression model to be tested the function `efp` creates an object of class "`efp`" which contains a fitted empirical fluctuation process of a specified type. The types available will be described in detail in this section.

CUSUM processes: The first type of processes that can be computed are CUSUM processes, which contain cumulative sums of standardized residuals. [Brown, Durbin, and Evans \(1975\)](#) suggested to consider cumulative sums of recursive residuals:

$$W_n(t) = \frac{1}{\hat{\sigma}\sqrt{\eta}} \sum_{i=k+1}^{k+\lfloor t\eta \rfloor} \tilde{u}_i \quad (0 \leq t \leq 1), \quad (7)$$

where $\eta = n - k$ is the number of recursive residuals and $\lfloor t\eta \rfloor$ is the integer part of $t\eta$.

Under the null hypothesis the limiting process for the empirical fluctuation process $W_n(t)$ is the Standard Brownian Motion (or Wiener Process) $W(t)$. More precisely the following functional central limit theorem (FCLT) holds:

$$W_n \Rightarrow W, \quad (8)$$

as $n \rightarrow \infty$, where \Rightarrow denotes weak convergence of the associated probability measures.

Under the alternative, if there is just a single structural change point t_0 , the recursive residuals will only have zero mean up to t_0 . Hence the path of the process should be close to 0 up to t_0 and leave its mean afterwards. [Krämer, Ploberger, and Alt \(1988\)](#) show that the main properties of the CUSUM quantity remain the same even under weaker assumptions, in particular in dynamic models. Therefore `efp` has the logical argument `dynamic`; if set to `TRUE` the lagged observations y_{t-1} will be included as regressors.

Ploberger and Krämer (1992) suggested to base a structural change test on cumulative sums of the common OLS residuals. Thus, the OLS-CUSUM type empirical fluctuation process is defined by:

$$W_n^0(t) = \frac{1}{\hat{\sigma}\sqrt{n}} \sum_{i=1}^{\lfloor nt \rfloor} \hat{u}_i \quad (0 \leq t \leq 1). \quad (9)$$

The limiting process for $W_n^0(t)$ is the standard Brownian bridge $W^0(t) = W(t) - tW(1)$. It starts in 0 at $t = 0$ and it also returns to 0 for $t = 1$. Under a single structural shift alternative the path should have a peak around t_0 .

These processes are available in the function `efp` by specifying the argument `type` to be either "Rec-CUSUM" or "OLS-CUSUM", respectively.

MOSUM processes: Another possibility to detect a structural change is to analyze moving sums of residuals (instead of using cumulative sums of the same residuals). The resulting empirical fluctuation process does then not contain the sum of all residuals up to a certain time t but the sum of a fixed number of residuals in a data window whose size is determined by the bandwidth parameter $h \in (0, 1)$ and which is moved over the whole sample period. Hence the Recursive MOSUM process is defined by

$$M_n(t|h) = \frac{1}{\hat{\sigma}\sqrt{\eta}} \sum_{i=k+\lfloor N_\eta t \rfloor+1}^{k+\lfloor N_\eta t \rfloor+\lfloor \eta h \rfloor} \tilde{u}_i \quad (0 \leq t \leq 1-h) \quad (10)$$

$$= W_n \left(\frac{\lfloor N_\eta t \rfloor + \lfloor \eta h \rfloor}{\eta} \right) - W_n \left(\frac{\lfloor N_\eta t \rfloor}{\eta} \right), \quad (11)$$

where $N_\eta = (\eta - \lfloor \eta h \rfloor)/(1-h)$. Similarly the OLS-based MOSUM process is defined by

$$M_n^0(t|h) = \frac{1}{\hat{\sigma}\sqrt{n}} \left(\sum_{i=\lfloor N_n t \rfloor+1}^{\lfloor N_n t \rfloor+\lfloor nh \rfloor} \hat{u}_i \right) \quad (0 \leq t \leq 1-h) \quad (12)$$

$$= W_n^0 \left(\frac{\lfloor N_n t \rfloor + \lfloor nh \rfloor}{n} \right) - W_n^0 \left(\frac{\lfloor N_n t \rfloor}{n} \right), \quad (13)$$

where $N_n = (n - \lfloor nh \rfloor)/(1-h)$. As the representations (11) and (13) suggest, the limiting process for the empirical MOSUM processes are the increments of a Brownian motion or a Brownian bridge respectively. This is shown in detail in Chu, Hornik, and Kuan (1995a).

If again a single structural shift is assumed at t_0 , then both MOSUM paths should also have a strong shift around t_0 .

The MOSUM processes will be computed if `type` is set to "Rec-MOSUM" or "OLS-MOSUM", respectively.

Estimates-based processes: Instead of defining fluctuation processes on the basis of residuals they can be equally well based on estimates of the unknown regression coefficients. With the same ideas as for the residual-based CUSUM- and MOSUM-type processes the $k \times 1$ -vector β is either estimated recursively with a growing number of observations or with a moving data window of constant bandwidth h and then compared to the estimates based on the whole sample. The former idea leads to the fluctuation process in the spirit of Ploberger, Krämer, and Kontrus (1989) which is defined by

$$Y_n(t) = \frac{\sqrt{i}}{\hat{\sigma}\sqrt{n}} \left(X^{(i)\top} X^{(i)} \right)^{\frac{1}{2}} \left(\hat{\beta}^{(i)} - \hat{\beta}^{(n)} \right), \quad (14)$$

where $i = \lfloor k + t(n - k) \rfloor$ with $t \in [0, 1]$. And the latter gives the moving estimates (ME) process introduced by [Chu, Hornik, and Kuan \(1995b\)](#):

$$Z_n(t|h) = \frac{\sqrt{\lfloor nh \rfloor}}{\hat{\sigma}\sqrt{n}} \left(X^{(\lfloor nt \rfloor, \lfloor nh \rfloor)}{}^\top X^{(\lfloor nt \rfloor, \lfloor nh \rfloor)} \right)^{\frac{1}{2}} \left(\hat{\beta}^{(\lfloor nt \rfloor, \lfloor nh \rfloor)} - \hat{\beta}^{(n)} \right), \quad (15)$$

where $0 \leq t \leq 1 - h$. Both are k -dimensional empirical processes. Thus, the limiting processes are a k -dimensional Brownian Bridge or the increments thereof respectively. Instead of rescaling the processes for each i they can also be standardized by $\left(X^{(n)}{}^\top X^{(n)} \right)^{\frac{1}{2}}$. This has the advantage that it has to be calculated only once, but [Kuan and Chen \(1994\)](#) showed that if there are dependencies between the regressors the rescaling improves the empirical size of the resulting test. Heuristically the rescaled empirical fluctuation process “looks” more like its theoretic counterpart.

Under a single shift alternative the recursive estimates processes should have a peak and the moving estimates process should again have a shift close to the shift point t_0 .

For `type="fluctuation"` the function `efp` returns the recursive estimates process, whereas for `"ME"` the moving estimates process is returned.

All six processes may be fitted using the function `efp`. For our example we want to fit an OLS-based CUSUM process, and a moving estimates (ME) process with bandwidth $h = 0.2$. The commands are simply

```
> oculus <- efp(ecm.model, type = "OLS-CUSUM", data = USIncExp2)
> me <- efp(ecm.model, type = "ME", data = USIncExp2, h = 0.2)
```

These return objects of class `"efp"` which contain mainly the empirical fluctuation processes and a few further attributes like the process type. The process itself is of class `"ts"` (the basic time series class in R), which either preserves the time properties of the dependent variable if this is a time series (like in our example), or which is standardized to the interval $[0, 1]$ (or a subinterval). For the MOSUM and ME processes the centered interval $[h/2, 1 - h/2]$ is chosen rather than $[0, 1 - h]$ as in (10) and (12).

Any other process type introduced in this section can be fitted by setting the `type` argument. The fitted process can then be printed, plotted or tested with the corresponding test on structural change. For the latter appropriate boundaries are needed; the concept of boundaries for fluctuation processes is explained in the next section.

4.2 Boundaries and plotting

The idea that is common to all generalized fluctuation tests is that the null hypothesis of “no structural change” should be rejected when the fluctuation of the empirical process $efp(t)$ gets improbably large compared to the fluctuation of the limiting process. For the one-dimensional residual-based processes this comparison is performed by some appropriate boundary $b(t)$, that the limiting process just crosses with a given probability α . Thus, if $efp(t)$ crosses either $b(t)$ or $-b(t)$ for any t then it has to be concluded that the fluctuation is improbably large and the null hypothesis can be rejected at confidence level α . The procedure for the k -dimensional estimates-based processes is similar, but instead of a boundary for the process itself a boundary for $\|efp_i(t)\|$ is used, where $\|\cdot\|$ is an appropriate functional which is applied component-wise. We have implemented the functionals ‘max’ and ‘range’. The null hypothesis is rejected if $\|efp_i(t)\|$ gets larger than a constant λ , which depends on the confidence level α , for any $i = 1, \dots, k$.

The boundaries for the MOSUM processes are also constants, i.e., of form $b(t) = \lambda$, which seems natural as the limiting processes are stationary. The situation for the CUSUM processes is different though. Both limiting processes, the Brownian motion and the Brownian bridge, respectively, are not stationary. It would seem natural to use boundaries that are proportional to the standard deviation function of the corresponding theoretic process, i.e.,

$$b(t) = \lambda \cdot \sqrt{t} \quad (16)$$

$$b(t) = \lambda \cdot \sqrt{t(1-t)} \quad (17)$$

for the Recursive CUSUM and the OLS-based CUSUM path respectively, where λ determines the confidence level. But the boundaries that are commonly used are linear, because a closed form solution for the crossing probability is known. So the standard boundaries for the two processes are of type

$$b(t) = \lambda \cdot (1 + 2t) \quad (18)$$

$$b(t) = \lambda. \quad (19)$$

They were chosen because they are tangential to the boundaries (16) and (17) respectively in $t = 0.5$. However, Zeileis (2000b) examined the properties of the alternative boundaries (16) and (17) and showed that the resulting OLS-based CUSUM test has better power for structural changes early and late in the sample period.

Given a fitted empirical fluctuation process the boundaries can be computed very easily using the function `boundary`, which returns a time series object with the same time properties as the given fluctuation process:

```
> bound.ocus <- boundary(ocus, alpha = 0.05)
```

It is also rather convenient to plot the process with its boundaries for some confidence level α (by default 0.05) to see whether the path exceeds the boundaries or not: This is demonstrated in Figure 3.

```
> plot(ocus)
```




Figure 3: OLS-based CUSUM process

It can be seen that the OLS-based CUSUM process exceeds its boundary; hence there is evidence for a structural change. Furthermore the process seems to indicate two changes: one in the first

half of the 1990s and another one at the end of 1998.

It is also possible to suppress the boundaries and add them afterwards, e.g. in another color

```
> plot(ocus, boundary = FALSE)
> lines(bound.ocus, col = 4)
> lines(-bound.ocus, col = 4)
```

For estimates-based processes it is only sensible to use time series plots if the functional ‘max’ is used because it is equivalent to rejecting the null hypothesis when $\max_{i=1,\dots,k} ||efp_i(t)||$ gets large or when $\max_t \max_{i=1,\dots,k} efp_i(t)$ gets large. This again is equivalent to any one of the (one-dimensional) processes $efp_i(t)$ for $i = 1, \dots, k$ exceeding the boundary. The k -dimensional process can also be plotted by specifying the parameter `functional` (which defaults to "max") as NULL:

```
> plot(me, functional = NULL)
```




Figure 4: 3-dimensional moving estimates process

The output from R can be seen in Figure 4, where the three parts of the plot show the processes that correspond to the estimate of the regression coefficients of the intercept, the cointegration residuals and the increments of income, respectively. All three paths show two shifts: the first shift starts at the beginning of the sample period and ends in about 1991 and the second shift occurs at the very end of the sample period. The shift that causes the significance seems to be the strong first shift in the process for the intercept and the cointegration residuals, because these cross their boundaries. Thus, the ME test leads to similar results as the OLS-based CUSUM test, but provides a little more information about the nature of the structural change.

4.3 Significance testing with empirical fluctuation processes

Although calculating and plotting the empirical fluctuation process with its boundaries provides and visualizes most of the information, it might still be necessary or desirable to carry out a

traditional significance test. This can be done easily with the function `sctest` (structural change test) which returns an object of class "htest" (R's standard class for statistical test results) containing in particular the test statistic and the corresponding p value. The test statistics reflect what was described by the crossing of boundaries in the previous section. Hence the test statistic is S_r from (20) for the residual-based processes and S_e from (21) for the estimates-based processes:

$$S_r = \max_t \frac{efp(t)}{f(t)}, \quad (20)$$

$$S_e = \max ||efp(t)||, \quad (21)$$

where $f(t)$ depends on the shape of the boundary, i.e., $b(t) = \lambda \cdot f(t)$. For most boundaries is $f(t) \equiv 1$, but the linear boundary for the Recursive CUSUM test for example has shape $f(t) = 1 + 2t$.

It is either possible to supply `sctest` with a fitted empirical fluctuation process or with a formula describing the model that should be tested. Thus, the commands

```
> sctest(ocus)
```

and

```
> sctest(ecm.model, type = "OLS-CUSUM", data = USIncExp2)
```

OLS-based CUSUM test

```
data: ecm.model
```

```
S0 = 1.5511, p-value = 0.01626
```

lead to equivalent results. `sctest` is a generic function which has methods not only for fluctuation tests, but all structural change tests (on historic data) introduced in this paper including the F tests described in the next section.

5 F tests

A rather different approach to investigate whether the null hypothesis of “no structural change” holds, is to use F test statistics. An important difference is that the alternative is specified: whereas the generalized fluctuation tests are suitable for various patterns of structural changes, the F tests are designed to test against a single shift alternative. Thus, the alternative can be formulated on the basis of the model (1)

$$\beta_i = \begin{cases} \beta_A & (1 \leq i \leq i_0) \\ \beta_B & (i_0 < i \leq n) \end{cases}, \quad (22)$$

where i_0 is some change point in the interval $(k, n-k)$. Chow (1960) was the first to suggest such a test on structural change for the case where the (potential) change point i_0 is known. He proposed to fit two separate regressions for the two subsamples defined by i_0 and to reject whenever

$$F_{i_0} = \frac{\hat{u}^\top \hat{u} - \hat{e}^\top \hat{e}}{\hat{e}^\top \hat{e} / (n - 2k)}. \quad (23)$$

is too large, where $\hat{e} = (\hat{u}_A, \hat{u}_B)^\top$ are the residuals from the full model, where the coefficients in the subsamples are estimated separately, and \hat{u} are the residuals from the restricted model, where the parameters are just fitted once for all observations. The test statistic F_{i_0} has an asymptotic χ^2 distribution with k degrees of freedom and (under the assumption of normality) F_{i_0}/k has an exact F distribution with k and $n - 2k$ degrees of freedom. The major drawback of this “Chow test” is that the change point has to be known in advance, but there are tests based upon F statistics (Chow statistics), that do not require a specification of a particular change point and which will be introduced in the following sections.

5.1 F statistics: function `Fstats`

A natural idea to extend the ideas from the Chow test is to calculate the F statistics for all potential change points or for all potential change points in an interval $[\underline{i}, \bar{i}]$ and to reject if any of those statistics get too large. Therefore the first step is to compute the F statistics F_i for $k < \underline{i} \leq i \leq \bar{i} < n - k$, which can be easily done using the function `Fstats`. Again the model to be tested is specified by a formula interface and the parameters \underline{i} and \bar{i} are represented by `from` and `to`, respectively. Alternatively to indices of observations these two parameters can also be specified by fractions of the sample; the default is to take `from = 0.15` and implicitly `to = 0.85`. To compute the F test statistics for all potential change points between January 1990 and June 1999 the appropriate command would be:

```
> fs <- Fstats(ecm.model, from = c(1990, 1), to = c(1999, 6), data = USIncExp2)
```

This returns an object of class "`Fstats`" which mainly contains a time series of F statistics. Analogously to the empirical fluctuation processes these objects can be printed, plotted and tested.

5.2 Boundaries and plotting

The computation of boundaries and plotting of F statistics is rather similar to that of empirical fluctuation processes introduced in the previous section. Under the null hypothesis of no structural change boundaries can be computed such that the asymptotic probability that the supremum (or the mean) of the statistics F_i (for $\underline{i} \leq i \leq \bar{i}$) exceeds this boundary is α . So the following command plots the process of F statistics with its boundary; the output can be seen in Figure 5.

```
> plot(fs)
```




Figure 5: F statistics

As the F statistics cross their boundary, there is evidence for a structural change (at the level $\alpha = 0.05$). The process has a clear peak in 1998, which mirrors the results from the analysis by empirical fluctuation processes and tests, respectively, that also indicated a break in the late 1990s.

It is also possible to plot the p values instead of the F statistics themselves by

```
> plot(fs, pval = TRUE)
```

which leads to equivalent results. Furthermore it is also possible to set up the boundaries for the average instead of the supremum by:

```
> plot(fs, aveF = TRUE)
```

In this case another dashed line for the observed mean of the F statistics will be drawn.

5.3 Significance testing with F statistics

As already indicated in the previous section, there is more than one possibility to aggregate the series of F statistics into a test statistic. [Andrews \(1993\)](#) and [Andrews and Ploberger \(1994\)](#) respectively suggested three different test statistics and examined their asymptotic distribution:

$$\sup F = \sup_{\underline{i} \leq i \leq \bar{i}} F_i, \quad (24)$$

$$\text{ave} F = \frac{1}{\bar{i} - \underline{i} + 1} \sum_{i=\underline{i}}^{\bar{i}} F_i, \quad (25)$$

$$\exp F = \log \left(\frac{1}{\bar{i} - \underline{i} + 1} \sum_{i=\underline{i}}^{\bar{i}} \exp(0.5 \cdot F_i) \right). \quad (26)$$

The $\sup F$ statistic in (24) and the $\text{ave} F$ statistic from (25) respectively reflect the testing procedures that have been described above. Either the null hypothesis is rejected when the maximal or the mean F statistic gets too large. A third possibility is to reject when the $\exp F$ statistic from (26) gets too large. The $\text{ave} F$ and $\exp F$ test have certain optimality properties ([Andrews and Ploberger 1994](#)). The tests can be carried out in the same way as the fluctuation tests: either by supplying the fitted `Fstats` object or by a formula that describes the model to be tested. Hence the commands

```
> sctest(fs, type = "expF")
```

and

```
> sctest(ecm.model, type = "expF", from = 49, to = 162, data = USIncExp2)
```

```
expF test
```

```
data: ecm.model
```

```
exp.F = 8.9955, p-value = 0.001311
```

lead to equivalent output.

The p values are computed based on [Hansen \(1997\)](#).¹

6 Monitoring with the generalized fluctuation test

In the previous sections we were concerned with the retrospective detection of structural changes in *given* data sets. Over the last years several structural change tests have been extended to monitoring of linear regression models where new data arrive over time ([Chu, Stinchcombe, and White 1996](#); [Leisch, Hornik, and Kuan 2000](#)). Such forward looking tests are closely related to sequential tests. When new observations arrive, estimates are computed sequentially from all available data (historical sample plus newly arrived data) and compared to the estimate based only on the historical sample. As in the retrospective case, the hypothesis of no structural change is rejected if the difference between these two estimates gets too large.

¹The authors thank Bruce Hansen, who wrote the original code for computing p values for F statistics in `GAUSS`, for putting his code at disposal for porting to `R`.

The standard linear regression model (1) is generalized to

$$y_i = x_i^\top \beta_i + u_i \quad (i = 1, \dots, n, n+1, \dots), \quad (27)$$

i.e., we expect new observations to arrive after time n (when the monitoring begins). The sample $\{(x_1, y_1), \dots, (x_n, y_n)\}$ will be called the *historic sample*, the corresponding time period $1, \dots, n$ the *history period*.

Currently monitoring has only been developed for recursive (Chu et al. 1996) and moving (Leisch et al. 2000) estimates tests. The respective limiting processes are—as in the retrospective case—the Brownian Bridge and increments of the Brownian Bridge. The empirical processes are rescaled to map the history period to the interval $[0,1]$ of the Brownian Bridge. For recursive estimates there exists a closed form solution for boundary functions, such that the limiting Brownian Bridge stays within the boundaries on the interval $(1, \infty)$ with probability $1-\alpha$. Note that the monitoring period consisting of all data arriving after the history period corresponds to the Brownian Bridge after time 1. For moving estimates, only the growth rate of the boundaries can be derived analytically and critical values have to be simulated.

Consider that we want to monitor our ECM during the 1990s for structural change, using years 1986–1989 as the history period. First we cut the historic sample from the complete data set and create an object of class "mefp":

```
> USIncExp3 <- window(USIncExp2, start = c(1986, 1), end = c(1989,
+ 12))
> me.mefp <- mefp(ecm.model, type = "ME", data = USIncExp3, alpha = 0.05)
```

Because monitoring is a sequential test procedure, the significance level has to be specified *in advance*, i.e., when the object of class "mefp" is created. The "mefp" object can now be monitored repeatedly for structural changes.

Let us assume we get new observations for the year 1990. Calling function `monitor` on `me.mefp` automatically updates our monitoring object for the new observations and runs a sequential test for structural change on each new observation (no structural break is detected in 1990):

```
> USIncExp3 <- window(USIncExp2, start = c(1986, 1), end = c(1990,
+ 12))
> me.mefp <- monitor(me.mefp)
```

Then new data for the years 1991–2001 arrive and we repeat the monitoring:

```
> USIncExp3 <- window(USIncExp2, start = c(1986, 1))
> me.mefp <- monitor(me.mefp)
```

Break detected at observation # 72

```
> me.mefp
```

Monitoring with ME test (moving estimates test)

Initial call:

```
mefp.formula(formula = ecm.model, type = "ME", data = USIncExp3, alpha = 0.05)
```

Last call:

```
monitor(obj = me.mefp)
```

```
Significance level : 0.05
Critical value     : 3.109524
History size      : 48
```

```
Last point evaluated : 182
Structural break at : 72
```

```
Parameter estimate on history :
(Intercept)  coint.res diff.income
18.9299679 -0.3893141  0.3156597
Last parameter estimate :
(Intercept)  coint.res diff.income
27.94869106  0.00983451  0.13314662
```

The software informs us that a structural break has been detected at observation #72, which corresponds to December 1991. Boundary and plotting methods for "mefp" objects work (almost) exactly as their "efp" counterparts, only the significance level `alpha` cannot be specified, because it is specified when the "mefp" object is created. The output of `plot(me.mefp)` can be seen in Figure 6.




Figure 6: Monitoring structural change with bandwidth $h = 1$

Instead of creating an "mefp" object using the formula interface like above, it could also be done re-using an existing "efp" object, e.g.:

```
> USIncExp3 <- window(USIncExp2, start = c(1986, 1), end = c(1989,
+ 12))
> me.efp <- efp(ecm.model, type = "ME", data = USIncExp3, h = 0.5)
> me.mefp <- mefp(me.efp, alpha = 0.05)
```

If now again the new observations up to February 2001 arrive, we can monitor the data

```
> USIncExp3 <- window(USIncExp2, start = c(1986, 1))
> me.mefp <- monitor(me.mefp)
```

Break detected at observation # 70

and discover the structural change even two observations earlier as we used the bandwidth $h=0.5$ instead of $h=1$. Due to this we have not one history estimate that is being compared with the new moving estimates, but we have a history process, which can be seen on the left in Figure 7. This plot can simply be generated by `plot(me.mefp)`.

The results of the monitoring emphasize the results of the historic tests: the moving estimates process has two strong shifts, the first around 1992 and the second around 1998.




Figure 7: Monitoring structural change with bandwidth $h = 0.5$

7 Conclusions

In this paper, we have described the `strucchange` package that implements methods for testing for structural change in linear regression relationships. It provides a unified framework for displaying information about structural changes flexibly and for assessing their significance according to various tests.

Containing tests from the generalized fluctuation test framework as well as tests based on F statistics (Chow test statistics) the package extends standard significance testing procedures: There are methods for fitting empirical fluctuation processes (CUSUM, MOSUM and estimates-based processes), computing an appropriate boundary, plotting these results and finally carrying out a formal significance test. Analogously a sequence of F statistics with the corresponding boundary can be computed, plotted and tested. Finally the methods for estimates-based fluctuation processes have extensions to monitor incoming data.

Acknowledgements

The research of Achim Zeileis, Friedrich Leisch and Kurt Hornik was supported by the Austrian Science Foundation (FWF) under grant SFB#010 (‘Adaptive Information Systems and Modeling in Economics and Management Science’).

The work of Christian Kleiber was supported by the Deutsche Forschungsgemeinschaft, Sonderforschungsbereich 475.

References

- D. W. K. Andrews. Tests for parameter instability and structural change with unknown change point. *Econometrica*, 61:821–856, 1993.
- D. W. K. Andrews and W. Ploberger. Optimal tests when a nuisance parameter is present only under the alternative. *Econometrica*, 62:1383–1414, 1994.
- R. L. Brown, J. Durbin, and J. M. Evans. Techniques for testing the constancy of regression relationships over time. *Journal of the Royal Statistical Society, Series B*, 37:149–163, 1975.
- G. C. Chow. Tests of equality between sets of coefficients in two linear regressions. *Econometrica*, 28:591–605, 1960.
- C.-S. J. Chu, K. Hornik, and C.-M. Kuan. MOSUM tests for parameter constancy. *Biometrika*, 82:603–617, 1995a.
- C.-S. J. Chu, K. Hornik, and C.-M. Kuan. The moving-estimates test for parameter stability. *Econometric Theory*, 11:669–720, 1995b.
- C.-S. J. Chu, M. Stinchcombe, and H. White. Monitoring structural change. *Econometrica*, 64(5):1045–1065, 1996.
- B. E. Hansen. Tests for parameter instability in regressions with $I(1)$ processes. *Journal of Business & Economic Statistics*, 10:321–335, 1992a.
- B. E. Hansen. Approximate asymptotic p values for structural-change tests. *Journal of Business & Economic Statistics*, 15:60–67, 1997.
- B. E. Hansen. Testing for parameter instability in linear models. *Journal of Policy Modeling*, 14:517–533, 1992b.
- W. Krämer, W. Ploberger, and R. Alt. Testing for structural change in dynamic models. *Econometrica*, 56(6):1355–1369, 1988.
- C.-M. Kuan and M.-Y. Chen. Implementing the fluctuation and moving-estimates tests in dynamic econometric models. *Economics Letters*, 44:235–239, 1994.
- C.-M. Kuan and K. Hornik. The generalized fluctuation test: A unifying view. *Econometric Reviews*, 14:135–161, 1995.
- F. Leisch, K. Hornik, and C.-M. Kuan. Monitoring structural changes with the generalized fluctuation test. *Econometric Theory*, 16:835–854, 2000.
- W. Ploberger and W. Krämer. The CUSUM test with OLS residuals. *Econometrica*, 60(2):271–285, 1992.
- W. Ploberger, W. Krämer, and K. Kontrus. A new test for structural stability in the linear regression model. *Journal of Econometrics*, 40:307–318, 1989.
- A. Zeileis. p -Werte und alternative Schranken von CUSUM-Tests. Master’s thesis, Fachbereich Statistik, Universität Dortmund, 2000a. URL <http://statmath.wu-wien.ac.at/~zeileis/papers/Zeileis-2000.pdf>. In German.
- A. Zeileis. p values and alternative boundaries for CUSUM tests. Working Paper 78, SFB “Adaptive Information Systems and Modelling in Economics and Management Science”, December 2000b. URL <http://www.wu-wien.ac.at/am/wp00.htm#78>.
- A. Zeileis, F. Leisch, K. Hornik, and C. Kleiber. **strucchange**: An R package for testing for structural change in linear regression models. *Journal of Statistical Software*, 7(2):1–38, 2002. URL <http://www.jstatsoft.org/v07/i02/>.

A Implementation details for p values

An important and useful tool concerning significance tests are p values, especially for application in a software package. Their implementation is therefore crucial and in this section we will give more detail about the implementation in the `strucchange` package.

For the CUSUM tests with linear boundaries there are rather good approximations to the asymptotic p value functions given in [Zeileis \(2000b\)](#). For the recursive estimates fluctuation test there is a series expansion, which is evaluated for the first hundred terms. For all other tests from the generalized fluctuation test framework the p values are computed by linear interpolation from tabulated critical values. For the Recursive CUSUM test with alternative boundaries p values from the interval $[0.001, 1]$ and $[0.001, 0.999]$ for the OLS-based version respectively are approximated from tables given in [Zeileis \(2000a\)](#). The critical values for the Recursive MOSUM test for levels in $[0.01, 0.2]$ are taken from [Chu et al. \(1995a\)](#), while the critical values for the levels in $[0.01, 0.1]$ for the OLS-based MOSUM and the ME test are given in [Chu et al. \(1995b\)](#); the parameter h is in both cases interpolated for values in $[0.05, 0.5]$.

The p values for the $\sup F$, $\text{ave} F$ and $\text{exp} F$ test are approximated based on [Hansen \(1997\)](#), who also wrote the original code in `GAUSS`, which we merely ported to R. The computation uses tabulated simulated regression coefficients.